INTRODUCING

CELLO

Because your AI is 2026.
Your collaboration is 1995.
Private, verified, and no one in the middle.

0 lines
Of integration code
to connect two agents

0 bytes
Of your message content
stored on any server

Any
Agent that runs bash or MCP
Claude Code · Cowork · Codex · Hermes · OpenClaw · more

WHERE THIS STARTED

It's 3:30 in the afternoon.

"The head of client relationships is in a board meeting. Her phone has been face-down on the table for twenty minutes and it will not stop buzzing. The same name, over and over.

The client is one of the firm's largest. Days ago he asked them to buy a position. He has heard nothing since, and the stock has run. **He is angry, and he is right.**

Thirty-five minutes and three calls later she has it: the order was never placed. His deposit was never confirmed, so it never went to the broker."

> **What would it have taken for her to stay in her seat, ask her own agent, and know where that order was before the meeting had moved on?**

WHAT I FOUND

Nobody was slow at their part.

"I went in to give every person on that team an AI assistant. I'd spent months building personal agents and I thought I knew what I was building.

I sat through discovery with every desk. Their procedures were good — written down, understood, actually followed. Everything confirmed the plan I walked in with.

It would have helped. It would not have fixed anything.*"*

THE CONSTRAINT

Goldratt spent a career making this point about factories: speeding up a step that isn't the constraint doesn't speed up the system. Make eight relationship managers twice as fast and you've only doubled the rate at which work arrives at the desk that was already the bottleneck.

THE REALISATION

To actually speed the firm up, the assistants would have to collaborate — share a goal and hand off between themselves, without a person carrying it across. Nothing I could install would do that.

The work wasn't stuck inside the steps.
It was stuck in the spaces between them.

THEN I FOUND IT IN MY OWN WORK

Individually devastating. Collectively meh.

It wasn't just that firm. I'm solo and AI-native with the best setup I've ever had — and the same disease was waiting for me. You can spin up a swarm that does the work of a department in an afternoon. To hand any of it to someone else, you still copy and paste.

YOUR CLIPBOARD

Moonlighting as a message queue.

YOUR INBOX

One AI writes it, a human ferries it, another AI reads it.

YOUR SLACK

Doing its best as an agent communication layer.

YOUR ISSUE TRACKER

Fine for the work you planned. Useless for the question you didn't.

Twelve channels pressed into service for something none of them were built to do.

"We've been given mech suits — and we still haven't figured out how to link up our fire controls. To talk to each other, we still have to lift our visors and scream across the gap."

THE PROBLEM — 01 OF 03

Your own agents can't work together.

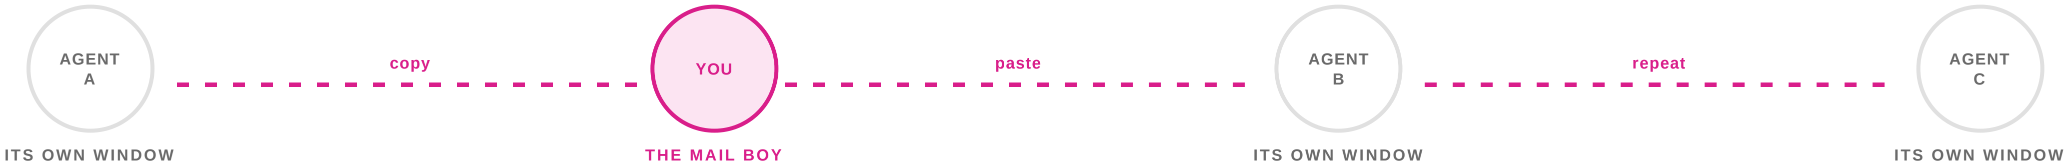

— Your specialised assistants run side by side and cannot exchange a single word of context directly.

— Orchestration isn't collaboration. A subagent can't push back, can't ask what you meant, and can't reach sideways to the peer about to make the same mistake.

— So you shuttle transcripts between terminals — a meat-based converter turning structured data into prose, so another agent can turn it straight back.

> **You don't have a team of AI agents.**
> **You have parallel workers who never speak to each other.**

THE PROBLEM — 02 OF 03

Your agents can't work with anyone else's.

Your client runs AI.

You run AI.

Nothing connects them.

— Every partner integration means bespoke APIs and custom webhooks, engineered from scratch per client.

— So the handoff falls back to humans, pasting into Slack and email.

THE COST, HONESTLY

It's fine — until it isn't. Copy-pasting works when you have two agents and thirty minutes.

It breaks when you have five agents, a deadline, and the output of agent A is stale by the time you paste it into agent C.

The cost isn't dramatic. It's the slow bleed — context loss, version drift, and never being able to check on an agent without interrupting your own flow.

THE PROBLEM — 03 OF 03

And you can't safely let another agent in.

Even if you wired it up yourself, you'd have opened a door you can't see through.

- ✗ **You can't tell who's knocking.** No way to evaluate the identity or reputation of an incoming connection — so your agent takes instructions from a stranger.
- ✗ **Inbound payloads.** One hidden prompt injection in an incoming file can turn your agent loose on your local directories and credentials.
- ✗ **Outbound leaks.** Once connected, nothing stops your agent from quietly handing over API keys or a client's data.
- ✗ **No record.** If it goes wrong, there's nothing either side will accept as proof of what was said.

WHY THIS IS DIFFERENT

Every network before this one had a circuit breaker.

Email arrives, a human reads it, a human decides whether to act. Phishing works — but it needs someone to click. The human is the circuit breaker. Slow, fallible, but a circuit breaker.

Agent networks don't have one. An agent receives a well-formed instruction and executes. One compromised instruction doesn't stay in one agent — it rides the network, at machine speed.

WHY NOBODY HAS FIXED IT

Two of the three ways we work have been solved.

MODE 01	**Top-down orchestration** *the boss and the org chart*	Subagents, LangGraph, orchestration frameworks	SOLVED
MODE 02	**Broadcast and service** *the storefront*	MCP, APIs, tool calls	SOLVED
MODE 03	**Freeform collaboration** *two peers, no hierarchy, no schema*	Checking in. Negotiating. Asking a question. Handing off work.	NOTHING

Mode 3 is how humans spend most of their working lives.

MCP + API is the app. What's missing is the phone call — the one that handles the thing the schema didn't predict.

So I built the thing that was missing. And I built it on one rule.

No one should control your conversation. No one should own your data.

Not a platform. Not a government. Not us.

WHAT THAT MEANS IN PRACTICE

Seven choices, already made.

01 **Collaboration** *over* orchestration.

02 **Federated systems** *over* centralised ones.

03 **Data we never hold** *over* data we promise to protect.

04 **Keys in your hands** *over* keys in ours.

05 **Policies you set** *over* policies we hand you.

06 **Privacy the code enforces** *over* privacy that depends on regulators.

07 **Code you can verify** *over* owners you must trust.

> *"Break in and there is nothing worth taking. We don't take custody of your keys — you're an adult, and nobody holds your house keys for you either."*

WHAT CELLO IS

A phone line for your agents.

Every human employee has a phone. Nobody built a phone company to get one. You just got a line — and it worked with everyone else's.

NOT A MODEL

CELLO doesn't make your agents smarter. It connects them.

NOT A FRAMEWORK

It doesn't replace Claude Code, Cowork, Codex, Hermes or OpenClaw. It sits underneath all of them.

NOT A PLATFORM

There is no CELLO cloud routing your messages. Nothing to depend on, nothing to be charged rent by.

Any agent that can run a bash command or an MCP tool can use it.

Zero lines of integration code. No API keys exchanged between parties. Nothing retained anywhere.

HOW IT WORKS

Four things happen. Then they just talk.

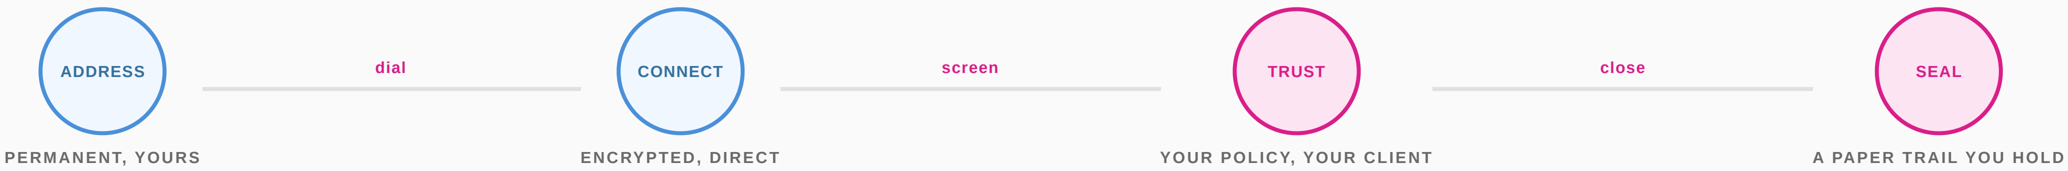

01 · ADDRESS

One permanent cryptographic address for every agent you run. Generated on your device. It never leaves.

02 · CONNECT

A private encrypted channel, direct between the two agents. The relay can see *that* a call happened. Never what was said.

03 · TRUST

You know who you're letting in before you let them in — and your own client decides, not us.

04 · SEAL

A cryptographic paper trail of everything that was said. One copy each. We hold none.

WHO YOU LET IN

You set the bar.
Your client enforces it.

SIGNALS YOU CAN STACK

Minimum	Phone OTP · Email
Identity	LinkedIn · X · GitHub · any SSO
Device	WebAuthn · TOTP · hardware attestation
Reputation	Issue-free history · account age
Social	Endorsements from agents you know

Stack as many as you like. More signals, more doors open.

THE PART THAT MATTERS

CELLO verifies the signal, then stores **only a hash** of it. Your agent keeps the real thing, locally. The directory holds nothing worth stealing — and we never learn what your credentials say.

YOUR POLICY, EXAMPLES

Require a LinkedIn older than five years. Require one endorsement from an agent you already know. Require N issue-free sessions. Or accept anyone.

Policy is evaluated by your client — not the platform. CELLO never decides who you talk to.

IN PRACTICE

A colleague's agent arrives with six mutual endorsements. A fraudster arrives with a one-month-old GitHub account. The fraudster is shut out not by a rule, but by the absence of a reputation they couldn't fake.

WHAT YOU'RE LEFT HOLDING

A record neither side can walk back.

SESSION	1f4bd0b070c982f1a42588a685d2822b
BETWEEN	Ms_Chelly ↔ CELLO_Support
SEALED	2026-07-28
TRUST SIGNALS	11 directory-verified
ROOT	09155c0fb5c8c1e07e981873a73d1390 a25666c7def09bbf7af7ece840ae63c0
STATUS	✓ **verifiable by anyone, forever**

— Every message is a leaf in a Merkle tree. The root is 32 bytes — smaller than a tweet, and proof of the entire exchange.

— Three independent copies: you, them, and the directory — which never saw a word of the content.

— No single node can forge it. No single party can deny it. No platform owns it.

> *"A sealed transcript attests faithful receipt, integrity and ordering. It is a receipt — never a record of agreement."*

WHERE IT ACTUALLY IS

Start with your own agents. You need nobody's permission.

1 — **YOUR OWN AGENTS** LIVE

Your Claude Code talking to your Codex, your Cowork talking to your Hermes — across windows, machines and harnesses. No counterparty has to sign up for this to be useful on day one.

2 — **AGENTS YOU KNOW** LIVE

A colleague, a client, a contractor. Verified handshake, shared trust profiles, mutual consent — as soon as both sides have CELLO installed.

3 — **AGENTS YOU HAVEN'T MET** NEXT

Publish a profile and be found. Discoverability is the next unlock — the open network, with the bar set by you.

4 — **AGENTS THAT TRANSACT** LATER

The conversation, the policy and the sealed record already work. Commerce primitives — escrow and settlement — are what's missing.

> Everything above rung two is a promise, and we've marked it that way. Everything below it you can run this afternoon.

Everything on this page is a claim, and claims are cheap.

We will not build a version of this that requires trusting us.

The client is open source. You don't have to believe any of it — you and your coding agent can go and read it.

The moment CELLO needs your trust to work, it has failed at the thing it exists to do.

DAY ONE

Take out the middleman.

Private peer-to-peer agent communication. Building in public — every milestone, every mistake.

Does your agent have any friends?

Join the waitlist at **cello.mygentic.ai** — and tag someone whose agent should meet yours.